

# More Data Structures

IS 665  
Yegin Genc

# Agenda

- Matrices
- Arrays
- Lists
- Dataframes
- Structures of structures

# Vector structures, starting with arrays

Many data structures in R are made by adding bells and whistles to vectors, so “vector structures”

A **matrix** in R is a collections of homogeneous elements arranged in 2 dimensions

```
matrix(1:15, nrow = 4)
```

	[,1]	[,2]	[,3]	[,4]
[1,]	1	5	9	13
[2,]	2	6	10	14
[3,]	3	7	11	15
[4,]	4	8	12	1

# Matrices

- A matrix is a vector with a `dim` attribute, i.e. an integer vector giving the number of rows and columns
- To create matrices use `matrix()`
- The functions `dim()`, `nrow()` and `ncol()` provide the attributes of the matrix
- Rows and columns can have names, `dimnames()`, `rownames()`, `colnames()`

# Matrices

```
factory <- matrix(c(40,1,60,3),nrow=2 )  
is.array(factory)
```

```
[1] TRUE
```

```
is.matrix(factory)
```

```
[1] TRUE
```

could also specify ncol, and/or byrow=TRUE to fill by rows.

Element-wise operations with the usual arithmetic and comparison operators (e.g., `factory/3`)

Compare whole matrices with `identical()` or `all.equal()`

# Matrix multiplication

Gets a special operator

```
six.sevens <- matrix(rep(7,6),ncol=3)  
six.sevens
```

```
      [,1] [,2] [,3]  
[1,]    7    7    7  
[2,]    7    7    7
```

```
factory %*% six.sevens # [2x2] * [2x3]
```

```
      [,1] [,2] [,3]  
[1,]  700  700  700  
[2,]   28   28   28
```

What happens if you try `six.sevens %*% factory`?

# Matrix operators

Transpose:

```
t(factory)
```

```
      [,1] [,2]  
[1,]  40   1  
[2,]  60   3
```

Determinant:

```
det(factory)
```

```
[1] 60
```

# Names in matrices

- We can name either rows or columns or both, with `rownames()` and `colnames()`
- These are just character vectors, and we use the same function to get and to set their values
- Names help us understand what we're working with
- Names can be used to coordinate different objects



```
rownames(factory) <- c("labor","steel")
colnames(factory) <- c("cars","trucks")
factory
```

```
      cars trucks
labor  40    60
steel  1     3
```

```
available <- c(1600,70)
names(available) <- c("labor","steel")
```

# Doing the same thing to each row or column

Take the mean: `rowMeans()`, `colMeans()`: input is matrix, output is vector. Also `rowSums()`, etc.

`summary()`: vector-style summary of column

```
colMeans(factory)
```

```
cars trucks  
20.5 31.5
```

```
summary(factory)
```

```
cars      trucks  
Min.   :1.00  Min.   :3.00  
1st Qu.:10.75 1st Qu.:17.25  
Median :20.50 Median :31.50  
Mean   :20.50 Mean   :31.50  
3rd Qu.:30.25 3rd Qu.:45.75  
Max.   :40.00 Max.   :60.00
```

# Extra

`apply()`, takes 3 arguments: the array or matrix, then 1 for rows and 2 for columns, then name of the function to apply to each

```
rowMeans(factory)
```

```
labor steel  
50 2
```

```
apply(factory,1,mean)
```

```
labor steel  
50 2
```

What would `apply(factory,1,sd)` do?

# Arrays

**arrays** are basically matrices in higher dimensions

```
x <- c(7, 8, 10, 45, 70, 80, 100, 250)
x.arr <- array(x, dim=c(2,2,2))
x.arr
```

```
., 1
     [,1] [,2]
[1,]   7  10
[2,]   8  45

., 2
     [,1] [,2]
[1,]  70 100
[2,]  80 250
```

dim says how many rows and columns; filled by columns

Can have 3, 4, . . .  $n$  dimensional arrays; dim is a length- $n$  vector

# Arrays cntd.

Some properties of the array:

```
dim(x.arr)
```

```
[1] 2 2 2
```

```
is.vector(x.arr)
```

```
[1] FALSE
```

```
is.array(x.arr)
```

```
[1] TRUE
```

# Arrays cntd.

```
typeof(x.arr)
```

```
[1] "double"
```

```
str(x.arr)
```

```
num [1:2, 1:2, 1:2] 7 8 10 45 70 80 100 250
```

```
attributes(x.arr)
```

```
$dim  
[1] 2 2 2
```

`typeof()` returns the type of the *elements*

`str()` gives the **structure**: here, a numeric array, with three dimensions, both indexed 1–2, and then the actual numbers

Exercise: try all these with `x`

# Accessing and operating on arrays

Can access a 2-D array either by pairs of indices or by the underlying vector:

```
x <- c(7, 8, 10, 45)  
x.arr <- array(x, dim=c(2,2))  
x.arr
```

```
      [,1] [,2]  
[1,]    7   10  
[2,]    8   45
```

```
x.arr[1,2]
```

```
[1] 10
```

```
x.arr[3]
```

```
[1] 10
```

# Accessing and operating on arrays

Omitting an index means “all of it”:

```
x.arr[c(1:2),2]
```

```
[1] 10 45
```

```
x.arr[,2]
```

```
[1] 10 45
```



# Functions on arrays

Using a vector-style function on a vector structure will go down to the underlying vector, *unless* the function is set up to handle arrays specially:

```
which(x.arr > 9)
```

```
[1] 3 4
```

# Functions on arrays

Many functions *do* preserve array structure:

```
y <- -x  
y.arr <- array(y,dim=c(2,2))  
y.arr + x.arr
```

```
      [,1] [,2]  
[1,]    0    0  
[2,]    0    0
```

Others specifically act on each row or column of the array separately:

```
rowSums(x.arr)
```

```
[1] 17 53
```

We will see a lot more of this idea

# Lists

Sequence of values, *not* necessarily all of the same type

```
my.distribution <- list("exponential",7,FALSE)  
my.distribution
```

```
[[1]]  
[1] "exponential"
```

```
[[2]]  
[1] 7
```

```
[[3]]  
[1] FALSE
```

Most of what you can do with vectors you can also do with lists

# Expanding and contracting lists

Add to lists with `c()` (also works with vectors):

```
my.distribution <- c(my.distribution,7)  
my.distribution
```

```
[[1]]  
[1] "exponential"
```

```
[[2]]  
[1] 7
```

```
[[3]]  
[1] FALSE
```

```
[[4]]  
[1] 7
```

Chop off the end of a list by setting the length to something smaller (also works with vectors):

```
length(my.distribution)
```

```
[1] 4
```

```
length(my.distribution) <- 3  
my.distribution
```

```
[[1]]  
[1] "exponential"
```

```
[[2]]  
[1] 7
```

```
[[3]]  
[1] FALSE
```

# Extra - Accessing pieces of lists

Can use `[ ]` as with vectors

or use `[[ ]]`, but only with a single index

`[[ ]]` drops names and structures, `[ ]` does not

```
is.character(my.distribution)
```

```
[1] FALSE
```

```
is.character(my.distribution[[1]])
```

```
[1] TRUE
```

```
my.distribution[[2]]^2
```

```
[1] 49
```

What happens if you try `my.distribution[2]^2`?

What happens if you try `[[ ]]` on a vector?

# Dataframes

Dataframe = the classic data table,  $n$  rows for cases,  $p$  columns for variables

Not just a matrix because *columns can have different types*

Many matrix functions also work for dataframes (rowSums(), summary(), apply())

but no matrix multiplication of dataframes, even if all columns are numeric

# Dataframes, Encore

- 2D tables of data
- Each case/unit is a row
- Each variable is a column
- Variables can be of any type (numbers, text, Booleans, ...)
- Both rows and columns can get names



# Creating an example dataframe

```
library(datasets)  
states <- data.frame(state.x77, abb=state.abb, region=state.region, division=state.division)
```

`data.frame()` is combining here a pre-existing matrix (`state.x77`), a vector of characters (`state.abb`), and two vectors of qualitative categorical variables (**factors**; `state.region`, `state.division`)

# Column names are preserved or guessed if not explicitly set

```
colnames(states)
```

```
[1] "Population" "Income"      "Illiteracy" "Life.Exp"    "Murder"  
[6] "HS.Grad"    "Frost"       "Area"        "abb"         "region"  
[11] "division"
```

```
states[1,]
```

```
      Population Income Illiteracy Life.Exp Murder HS.Grad Frost Area  
Alabama   3615 3624      2.1 69.05 15.1 41.3 20 50708  
      abb region      division  
Alabama AL South East South Central
```

# Dataframe access

- By row and column index

```
states[49,3]
```

```
[1] 0.7
```

- By row and column names

```
states["Wisconsin","Illiteracy"]
```

```
[1] 0.7
```

# Dataframe access (cont'd)

- All of a row:

```
states["Wisconsin",]
```

	Population	Income	Illiteracy	Life.Exp	Murder	HS.Grad	Frost	Area
Wisconsin	4589	4468	0.7	72.48	3	54.5	149	54464
abb	region	division						
Wisconsin	WI	North Central	East	North Central				

Exercise: what class is `states["Wisconsin",]`?

# Dataframe access (cont'd.)

- All of a column:

```
head(states[,3])
```

```
[1] 2.1 1.5 1.8 1.9 1.1 0.7
```

```
head(states["literacy"])
```

```
[1] 2.1 1.5 1.8 1.9 1.1 0.7
```

```
head(states$literacy)
```

```
[1] 2.1 1.5 1.8 1.9 1.1 0.7
```

# Dataframe access (cont'd.)

- Rows matching a condition:

```
states[states$division=="New England", "Illiteracy"]
```

```
[1] 1.1 0.7 1.1 0.7 1.3 0.6
```

```
states[states$region=="South", "Illiteracy"]
```

```
[1] 2.1 1.9 0.9 1.3 2.0 1.6 2.8 0.9 2.4 1.8 1.1 2.3 1.7 2.2 1.4 1.4
```

# Replacing values

Parts or all of the dataframe can be assigned to:

```
summary(states$HS.Grad)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
37.80	48.05	53.25	53.11	59.15	67.30

```
states$HS.Grad <- states$HS.Grad/100  
summary(states$HS.Grad)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.3780	0.4805	0.5325	0.5311	0.5915	0.6730

```
states$HS.Grad <- 100*states$HS.Grad
```

# Adding rows and columns

We can add rows or columns to an array or data-frame with `rbind()` and `cbind()`, but be careful about forced type conversions

```
a.data.frame()  
rbind(a.data.frame,list(v1=-3,v2=-5,logicals=TRUE))  
rbind(a.data.frame,c(3,4,6))
```

\*Internally, a dataframe is basically a list of vectors



# Summary

- Matrices act like you'd hope they would
- Arrays add multi-dimensional structure to vectors
- Lists let us combine different types of data
- Dataframes are hybrids of matrices and lists, for classic tabular data

# References

- <http://www.stat.cmu.edu/~cshalizi/statcomp/>
- <https://www.r-project.org/>
- <https://www.rstudio.com/>