

Introduction to R - Basics of Data

IS 665

Agenda

- Course overview and mechanics
- Built-in data types
- Built-in functions and operators
- First data structures: Vectors and arrays

Why good analysts learn to program

*(and not rely on off-the-shelf products)

- *Independence*: Otherwise, you rely on someone else having given you exactly the right tool (which usually never exists)
- *Honesty*: Otherwise, you end up distorting your problem to match the tools you have
- *Clarity*: Making your method something a machine can do disciplines your thinking and makes it public; that's science

How this part of the class will work

- No programming knowledge presumed (although would help)
- Some math and stats. knowledge presumed (will be reviewed)
- General programming mixed with data-manipulation and data mining applications
- Class will be *very* cumulative

Mechanics

- Lectures: concepts, methods, examples
- In-class Assignments to try stuff out and get fast feedback
- HW weekly to do longer and more complex things
- Projects:
 - Project 1: Exploratory Data Analysis
 - Project 2: Predictive (or Prescriptive Analysis)

Functional programming in a nutshell:

2 sorts of things (**objects**): **data** and **functions**

- **Data**: things like 7, “seven”, 7.000, the matrix $\begin{bmatrix} 7 & 7 & 7 \\ 7 & 7 & 7 \end{bmatrix}$
- **Functions**: things like `log`, `+` (two arguments), `<` (two), `mod` (two), `mean` (one)

A function is a machine which turns input objects (**arguments**) into an output object (**return value**), possibly with **side effects**, according to a definite rule

Before functions, data

Different kinds of data object

All data is represented in binary format, by **bits** (TRUE/FALSE, YES/NO, 1/0)

- **Booleans** Direct binary values: TRUE or FALSE in R
- **Integers**: whole numbers (positive, negative or zero), represented by a fixed-length block of bits
- **Characters** fixed-length blocks of bits, with special coding; **strings** = sequences of characters
- **Floating point numbers**: a fraction (with a finite number of bits) times an exponent, like 1.87×10^6 , but in binary form
- **Missing or ill-defined values**: NA, NaN, etc.

R as a calculator - Operators

```
7+5
```

```
[1] 12
```

```
7-5
```

```
[1] 2
```

```
7*5
```

```
[1] 35
```

```
7^5
```

```
[1] 16807
```


Operators cont'd.

```
7/5
```

```
[1] 1.4
```

```
7 %% 5
```

```
[1] 2
```

```
7 %/% 5
```

```
[1] 1
```

Operators cont'd.

Comparisons are also binary operators; they take two objects, like numbers, and give a Boolean

```
7 > 5
```

```
[1] TRUE
```

```
7 < 5
```

```
[1] FALSE
```

```
7 >= 7
```

```
[1] TRUE
```

```
7 <= 5
```

```
[1] FALSE
```

Operators cont'd.

`7 == 5`

[1] FALSE

`7 != 5`

[1] TRUE

Boolean operators

Basically “and” and “or”:

```
(5 > 7) & (6*7 == 42)
```

```
[1] FALSE
```

```
(5 > 7) | (6*7 == 42)
```

```
[1] TRUE
```

More types

`typeof()` function returns the type

`is.foo()` functions return Booleans for whether the argument is of type *foo*

`as.foo()` (tries to) “cast” its argument to type *foo* — to translate it sensibly into a *foo*-type value

More types

```
typeof(7)
```

```
[1] "double"
```

```
is.numeric(7)
```

```
[1] TRUE
```

```
is.na(7)
```

```
[1] FALSE
```

```
is.na(7/0)
```

```
[1] FALSE
```

```
is.na(0/0)
```

```
[1] TRUE
```

Why is 7/0 not NA, but 0/0 is?

More types

```
is.character(7)
```

```
[1] FALSE
```

```
is.character("7")
```

```
[1] TRUE
```

```
is.character("seven")
```

```
[1] TRUE
```

```
is.na("seven")
```

```
[1] FALSE
```

More types

```
as.character(5/6)
```

```
[1] "0.833333333333333"
```

```
as.numeric(as.character(5/6))
```

```
[1] 0.8333333
```

```
6*as.numeric(as.character(5/6))
```

```
[1] 5
```

```
5/6 == as.numeric(as.character(5/6))
```

```
[1] FALSE
```

(why is that last FALSE?)

Data can have names

We can give names to data objects; these give us **variables**

A few variables are built in:

```
pi
```

```
[1] 3.141593
```

Variables can be arguments to functions or operators, just like constants:

```
pi*10
```

```
[1] 31.41593
```

```
cos(pi)
```

```
[1] -1
```

Assignment operators

Most variables are created with the **assignment operator**, <- or =

```
approx.pi <- 22/7  
approx.pi
```

```
[1] 3.142857
```

```
diameter.in.cubits = 10  
approx.pi*diameter.in.cubits
```

```
[1] 31.42857
```

Assignment operators

The assignment operator also changes values:

```
circumference.in.cubits <- approx.pi*diameter.in.cubits  
circumference.in.cubits
```

```
[1] 31.42857
```

```
circumference.in.cubits <- 30  
circumference.in.cubits
```

```
[1] 30
```

Important Tips

- Using names and variables makes code: easier to design, easier to debug, less prone to bugs, easier to improve, and easier for others to read
- Avoid “magic constants”; use named variables
- Named variables are a first step towards **abstraction**

The workspace

What names have you defined values for?

```
ls()
```

```
[1] "approx.pi"      "circumference.in.cubits"  
[3] "diameter.in.cubits"
```

```
objects()
```

```
[1] "approx.pi"      "circumference.in.cubits"  
[3] "diameter.in.cubits"
```

Getting rid of variables:

```
rm("circumference.in.cubits")  
ls()
```

```
[1] "approx.pi"      "diameter.in.cubits"
```

First data structure: vectors

Group related data values into one object, a **data structure**

A **vector** is a sequence of values, all of the same type

```
x <- c(7, 8, 10, 45)  
x
```

```
[1] 7 8 10 45
```

```
is.vector(x)
```

```
[1] TRUE
```

c() function returns a vector containing all its arguments in order

x[1] is the first element, x[4] is the 4th element

x[-4] is a vector containing all but the fourth element

Vectors cont'd.

`vector(length=6)` returns an empty vector of length 6; helpful for filling things up later

```
weekly.hours <- vector(length=5)  
weekly.hours[5] <- 8
```

Vector arithmetic

Operators apply to vectors “pairwise” or “elementwise”:

```
y <- c(-7, -8, -10, -45)  
x+y
```

```
[1] 0 0 0 0
```

```
x*y
```

```
[1] -49 -64 -100 -2025
```


Recycling

Recycling repeats elements in shorter vector when combined with longer

```
x + c(-7,-8)
```

```
[1] 0 0 3 37
```

```
x^c(1,0,-1,0.5)
```

```
[1] 7.000000 1.000000 0.100000 6.708204
```

Single numbers are vectors of length 1 for purposes of recycling:

```
2*x
```

```
[1] 14 16 20 90
```

Comparison functions

Can also do pairwise comparisons:

```
x > 9
```

```
[1] FALSE FALSE TRUE TRUE
```

Note: returns Boolean vector

Boolean operators work elementwise:

```
(x > 9) & (x < 20)
```

```
[1] FALSE FALSE TRUE FALSE
```

Comparison functions

To compare whole vectors, best to use `identical()` or `all.equal()`:

```
x == -y
```

```
[1] TRUE TRUE TRUE TRUE
```

```
identical(x,-y)
```

```
[1] TRUE
```

```
identical(c(0.5-0.3,0.3-0.1),c(0.3-0.1,0.5-0.3))
```

```
[1] FALSE
```

```
all.equal(c(0.5-0.3,0.3-0.1),c(0.3-0.1,0.5-0.3))
```

```
[1] TRUE
```

Functions on vectors

Lots of functions take vectors as arguments:

- `mean()`, `median()`, `sd()`, `var()`, `max()`, `min()`, `length()`, `sum()`:
return single numbers
- `sort()` returns a new vector
- `hist()` takes a vector of numbers and produces a histogram, a highly structured object, with the side-effect of making a plot
- Similarly `ecdf()` produces a cumulative-density-function object
- `summary()` gives a five-number summary of numerical vectors
- `any()` and `all()` are useful on Boolean vectors

Addressing vectors

Vector of indices:

```
x[2];x[4]
```

```
[1] 8
```

```
[1] 45
```

```
x[c(2,4)]
```

```
[1] 8 45
```

Vector of negative indices

```
x[c(-1,-3)]
```

```
[1] 8 45
```

(why that, and not 8 10?)

Addressing Vectors cont'd.

Boolean vector:

```
x[x>9]
```

```
[1] 10 45
```

```
y[x>9]
```

```
[1] -10 -45
```

which() turns a Boolean vector in vector of TRUE indices:

```
places <- which(x > 9)  
places
```

```
[1] 3 4
```

```
y[places]
```

```
[1] -10 -45
```

Named components

You can give names to elements or components of vectors

```
names(x) <- c("v1", "v2", "v3", "fred")  
names(x)
```

```
[1] "v1" "v2" "v3" "fred"
```

```
x[c("fred", "v1")]
```

```
fred v1  
45  7
```

note the labels in what R prints; not actually part of the value

Named components

names(x) is just another vector (of characters):

```
names(y) <- names(x)  
sort(names(x))
```

```
[1] "fred" "v1"  "v2"  "v3"
```

```
which(names(x)=="fred")
```

```
[1] 4
```


Take-Aways

- We write programs by composing functions to manipulate data
- The basic data types let us represent Booleans, numbers, and characters
- Data structure let us group related values together
- Vectors let us group values of the same type
- Use variables rather a profusion of magic constants
- Name components of structures to make data more meaningful